

```
In [1]: # Import pandas
```

```
In [2]: # Create a series of three different colours
```

```
In [3]: # View the series of different colours
```

```
In [4]: # Create a series of three different car types and view it
```

```
In [5]: # Combine the Series of cars and colours into a DataFrame
```

```
In [6]: # Import "../data/car-sales.csv" and turn it into a DataFrame
```

**Note:** Since you've imported `../data/car-sales.csv` as a `DataFrame`, we'll now refer to this `DataFrame` as 'the car sales `DataFrame`'.

```
In [7]: # Export the DataFrame you created to a .csv file
```

```
In [8]: # Find the different datatypes of the car data DataFrame
```

```
In [9]: # Describe your current car sales DataFrame using describe()
```

```
In [10]: # Get information about your DataFrame using info()
```

What does it show you?

```
In [11]: # Create a Series of different numbers and find the mean of them
```

```
In [12]: # Create a Series of different numbers and find the sum of them
```

```
In [13]: # List out all the column names of the car sales DataFrame
```

```
In [14]: # Find the length of the car sales DataFrame
```

```
In [15]: # Show the first 5 rows of the car sales DataFrame
```

```
In [16]: # Show the first 7 rows of the car sales DataFrame
```

```
In [17]: # Show the bottom 5 rows of the car sales DataFrame
```

```
In [18]: # Use .loc to select the row at index 3 of the car sales DataFrame
```

```
In [19]: # Use .iloc to select the row at position 3 of the car sales DataFrame
```

Notice how they're the same? Why do you think this is?

Check the pandas documentation for [.loc](#) and [.iloc](#). Think about a different situation each could be used for and try them out.

```
In [20]: # Select the "Odometer (KM)" column from the car sales DataFrame
```

```
In [21]: # Find the mean of the "Odometer (KM)" column in the car sales DataFrame
```

```
In [22]: # Select the rows with over 100,000 kilometers on the Odometer
```

```
In [23]: # Create a crosstab of the Make and Doors columns
```

```
In [24]: # Group columns of the car sales DataFrame by the Make column and find the average
```

```
In [25]: # Import Matplotlib and create a plot of the Odometer column  
# Don't forget to use %matplotlib inline
```

```
In [26]: # Create a histogram of the Odometer column using hist()
```

```
In [27]: # Try to plot the Price column using plot()
```

Why didn't it work? Can you think of a solution?

You might want to search for "how to convert a pandas string column to numbers".

And if you're still stuck, check out this [Stack Overflow question and answer on turning a price column into integers](#).

See how you can provide the example code there to the problem here.

```
In [28]: # Remove the punctuation from price column
```

```
In [29]: # Check the changes to the price column
```

```
In [30]: # Remove the two extra zeros at the end of the price column
```

```
In [31]: # Check the changes to the Price column
```

```
In [32]: # Change the datatype of the Price column to integers
```

```
In [33]: # Lower the strings of the Make column
```

If you check the car sales DataFrame, you'll notice the Make column hasn't been lowered.

How could you make these changes permanent?

Try it out.

```
In [34]: # Make lowering the case of the Make column permanent
```

```
In [35]: # Check the car sales DataFrame
```

Notice how the Make column stays lowered after reassigning.

Now let's deal with missing data.

```
In [36]: # Import the car sales DataFrame with missing data ("../data/car-sales-missing-data.csv")
```

```
# Check out the new DataFrame
```

Notice the missing values are represented as NaN in pandas DataFrames.

Let's try fill them.

```
In [37]: # Fill the Odometer column missing values with the mean of the column inplace
```

```
In [38]: # View the car sales missing DataFrame and verify the changes
```

```
In [39]: # Remove the rest of the missing data inplace
```

```
In [40]: # Verify the missing values are removed by viewing the DataFrame
```

We'll now start to add columns to our DataFrame.

```
In [41]: # Create a "Seats" column where every row has a value of 5
```

```
In [42]: # Create a column called "Engine Size" with random values between 1.3 and 4.5  
# Remember: If you're doing it from a Python list, the list has to be the same length  
# as the DataFrame
```

```
In [43]: # Create a column which represents the price of a car per kilometer  
# Then view the DataFrame
```

```
In [44]: # Remove the last column you added using .drop()
```

```
In [45]: # Shuffle the DataFrame using sample() with the frac parameter set to 1  
# Save the the shuffled DataFrame to a new variable
```

Notice how the index numbers get moved around. The `sample()` function is a great way to get random samples from your DataFrame. It's also another great way to shuffle the rows by setting `frac=1`.

```
In [46]: # Reset the indexes of the shuffled DataFrame
```

Notice the index numbers have been changed to have order (start from 0).

```
In [47]: # Change the Odometer values from kilometers to miles using a Lambda function  
# Then view the DataFrame
```

```
In [48]: # Change the title of the Odometer (KM) to represent miles instead of kilometers
```